

Intelligent Automotive Design using Agentic AI

Course Name: Agentic AI

Institution Name: Medicaps University – Datagami Skill Based Course

Student Name(s) & Enrolment Number(s):

Sr no	Student Name	Enrolment Number
1	KANCHI TIWARI	EN22CS304032
2	ARYAN PAWAR	EN22CS304017
3	AMAN RAI	EN22CS304009
4	AAKRATI JAIN	EN22CS304001
5	PRATHAM BHAWAR	EN22CS304047
6	MOHAMMAD FARAZ ABBASI	EN22CS303033

Group Name: Group 10D11

Project Number: AAI-36

University Mentor Name: Prof. Nishant Shrivastav

Academic Year: 2026

Problem Statement & Objectives

1. Problem Statement
2. Project Objectives
3. Scope of the Project

Proposed Solution

1. Key features *(Just mention key features here no need to go into details)*
2. Overall Architecture / Workflow
3. Tools & Technologies Used *(If applicable)*

Results & Output

Add the below details here:

1. Screenshots / outputs
2. Reports / dashboards / models
3. Key outcomes

Conclusion

Mention a brief conclusion about your project summing up everything you have worked on and the key learning.

Future Scope & Enhancements

Problem Statement & Objective

1. Problem Statement

In the rapidly evolving automotive industry, accessing comprehensive and reliable information about vehicles—such as technical specifications, performance metrics, features, and comparisons—often requires collecting data from multiple heterogeneous sources. This information is typically unstructured and scattered across different platforms, making it difficult to interpret and organize into a meaningful format. As a result, users must invest significant time and effort in manually searching, verifying, and structuring the required data.

Traditional approaches rely either on manual processes or on monolithic artificial intelligence systems that attempt to handle data retrieval, processing, and presentation within a single pipeline. These systems lack modularity and flexibility, making them difficult to scale and maintain. Moreover, the absence of clear task separation often leads to inconsistent outputs and reduced efficiency.

With the advancement of Agentic Artificial Intelligence, a new paradigm has emerged in which multiple intelligent agents collaborate to solve complex problems by dividing tasks based on specialization. This approach enables more efficient task execution, improved output quality, and enhanced system scalability. However, implementing such systems introduces challenges related to effective task distribution, seamless inter-agent communication, context preservation, and maintaining consistency across multiple stages of processing.

This project addresses these challenges by designing and implementing a collaborative multi-agent automotive system in which specialized agents perform distinct roles. A Researcher Agent is responsible for retrieving relevant automotive data from structured or semi-structured sources, while a Writer Agent transforms this data into coherent, structured, and user-friendly outputs. The primary goal is to develop a system in which these agents can collaborate effectively through well-defined handoff protocols, enabling the transformation of a simple user query into a high-quality automotive report.

2. Project Objectives

The primary objective of this project is to design and implement a collaborative Agentic AI system that automates the process of automotive data retrieval and structured content generation. The project aims to develop a modular multi-agent architecture in which tasks are decomposed and assigned to specialized agents, thereby improving system efficiency, scalability, and maintainability.

The system is designed to include a Researcher Agent capable of extracting relevant automotive data, identifying key attributes such as engine specifications, performance metrics, and pricing, and structuring this information into a machine-readable format. In parallel, a Writer Agent is developed to interpret the structured data and generate well-organized, coherent, and readable

outputs in the form of reports, summaries, or comparison tables.

Another key objective is to establish a robust inter-agent communication mechanism using standardized handoff protocols, ensuring smooth data transfer and consistency in input and output formats. The project also focuses on implementing state management techniques to maintain context and track intermediate results throughout the workflow, thereby enabling continuity and accuracy in multi-step processing.

Furthermore, the system incorporates an orchestration layer responsible for managing the overall workflow, assigning tasks to appropriate agents, and ensuring efficient execution. Ultimately, the project aims to generate high-quality structured outputs such as car specification sheets, comparative analyses, and descriptive summaries while demonstrating the practical applicability of agent-based AI systems in real-world scenarios.

3. Scope of the Project

The scope of this project focuses on the design, development, and demonstration of a collaborative Agentic AI system for automotive information processing. The system is designed to accept natural language queries from users related to automotive data, such as vehicle specifications or comparisons, and process them through a structured multi-agent pipeline.

The workflow involves an orchestrator that manages the interaction between agents, a Researcher Agent that retrieves relevant data from predefined sources or knowledge bases, and a Writer Agent that converts the retrieved data into structured and human-readable outputs. The system is capable of generating outputs in various formats, including descriptive summaries, bullet-style specifications, and comparison reports.

The implementation includes the development of backend logic using Python and large language models, along with a basic frontend interface for user interaction. It also incorporates mechanisms for inter-agent communication, workflow orchestration, and state management to ensure efficient processing.

However, the project is limited to a controlled or local environment and does not include large-scale deployment or integration with live external APIs. Features such as multi-user authentication, real-time data streaming, and advanced personalization are beyond the current scope. The primary emphasis of the project is on demonstrating the architecture, agent collaboration, and workflow design of an Agentic AI system rather than building a fully production-ready application.

Proposed Solution

The proposed solution is a **Collaborative Agentic AI-based Automotive Information System** designed to automate the process of retrieving, processing, and presenting automotive data using multiple specialized agents. The system follows an agent-based architecture where different components are assigned specific roles, enabling efficient task execution, better modularity, and improved output quality. Unlike traditional monolithic AI systems, this approach divides the workflow into smaller, manageable tasks handled by independent agents, thereby enhancing scalability and maintainability.

The system allows users to input queries in natural language related to automotive information, such as vehicle specifications or comparisons. These queries are processed through a structured pipeline involving multiple agents that collaborate to generate meaningful and well-organized outputs. The system ensures accurate interpretation of user queries, efficient retrieval of relevant data, and transformation of that data into structured and readable formats.

At its core, the system accepts user queries in natural language and transforms them into structured automotive outputs through a coordinated workflow. This transformation involves multiple stages, including query understanding, data retrieval, data structuring, content synthesis, and final presentation. The system ensures that each stage is handled by a specialized component, enabling better accuracy and modularity. The use of agent-based architecture also allows the system to be easily extended in the future by adding new agents without disrupting existing functionality.

1. Key Features

The system is built with a range of advanced features that collectively enhance its performance, usability, and scalability. One of the primary features is the implementation of a multi-agent architecture, where different agents are assigned clearly defined roles, enabling specialization and reducing the cognitive load on any single component. This design improves both efficiency and maintainability, as each agent can be independently optimized or upgraded.

Another important feature is the support for natural language interaction, which allows users to input queries in a simple and intuitive manner without requiring technical knowledge or structured commands. The system is capable of interpreting these queries and converting them into actionable tasks. The inclusion of a Researcher Agent enables efficient retrieval of relevant automotive data, while the Writer Agent ensures that this data is transformed into structured, readable, and well-formatted outputs suitable for end users.

The system also incorporates a centralized orchestration mechanism that acts as a controller, managing the flow of information and coordinating interactions between agents. This orchestration ensures that tasks are executed in the correct sequence and that dependencies

between agents are properly handled. A standardized communication protocol is implemented to facilitate seamless data transfer between agents, ensuring consistency in data formats and reducing the chances of errors.

Additionally, the system includes a robust state management mechanism that tracks intermediate results and maintains context throughout the workflow. This is particularly important in multi-step processes where outputs from one stage serve as inputs to the next. The system is also capable of generating multiple types of structured outputs, including detailed car specifications, descriptive summaries, and comparative reports, making it versatile and applicable to a variety of use cases.

2. Overall Architecture / Workflow

The Collaborative Automotive Agent system is designed using a layered and modular architecture that ensures efficient interaction between user input, backend processing, and intelligent agent execution. The architecture separates responsibilities into distinct layers, enabling scalability, maintainability, and smooth data flow across the system. The overall workflow transforms a simple user query into a structured and high-quality automotive output through coordinated multi-agent collaboration.

Layered Architecture

The system architecture is divided into multiple layers, each responsible for a specific function within the workflow. This layered approach ensures clear separation of concerns and allows independent development and scaling of system components.

- **Presentation Layer (Frontend)**

The Presentation Layer serves as the entry point of the system and is responsible for handling user interaction. It is developed using technologies such as HTML, CSS, and JavaScript to provide a simple and intuitive interface.

This layer allows users to input automotive-related queries in natural language, such as requesting vehicle specifications or comparisons between different models. Once the user submits a query, the frontend validates the input and sends it to the backend using API requests. The frontend also includes components for displaying outputs in a structured manner, such as formatted text, specifications, or comparison results. Additionally, it may include user experience enhancements such as loading indicators and response messages to inform users about processing stages.

- **Application Layer (Backend)**

The Application Layer acts as the core processing unit of the system and is implemented using Python with frameworks such as Flask or FastAPI. This layer is responsible for handling incoming requests, managing workflow execution, and coordinating interactions between agents. At the heart of this layer is the **Orchestrator**, which functions as a controller responsible for task management and agent coordination. When a request is received, the orchestrator analyzes the user query to determine the required task, such as retrieving specifications or generating

comparisons. It then decomposes the task into smaller sub-tasks and assigns them to the appropriate agents.

This layer also includes a prompt processing module that refines and structures the input query to improve the performance of downstream agents. The backend ensures that all components communicate efficiently and that the workflow is executed in the correct sequence.

• **Agent Services Layer (AI Agent Layer)**

The Agent Services Layer is responsible for executing the core intelligence of the system through specialized agents. It consists of two primary agents: the Researcher Agent and the Writer Agent.

The Researcher Agent is responsible for retrieving relevant automotive data based on the user query. It processes the input, identifies key entities such as vehicle names and attributes, and extracts structured information such as specifications, performance metrics, and features from available data sources. The output of this agent is a structured dataset, typically in JSON format, which ensures consistency and ease of processing.

The Writer Agent is responsible for transforming the structured data into a coherent and user-friendly format. It interprets the data provided by the Researcher Agent and generates well-organized outputs such as descriptive summaries, specification sheets, or comparison reports. The Writer Agent ensures that the output is readable, logically structured, and aligned with the user's request.

The separation of responsibilities between these agents enables better task specialization, improved output quality, and easier system extensibility.

• **State Management Layer**

The State Management Layer plays a critical role in maintaining context and tracking the progress of the workflow. Since the system involves multiple stages of processing, it is essential to store intermediate outputs and ensure that data flows correctly between agents.

This layer maintains information such as the original user query, intermediate results generated by the Researcher Agent, and the final output produced by the Writer Agent. By preserving this state, the system ensures consistency and allows for accurate processing across multiple steps. It also facilitates debugging and monitoring by providing visibility into each stage of execution.

• **Communication Layer**

The Communication Layer ensures seamless interaction between different components of the system. It is based on RESTful APIs and uses JSON as the standard data exchange format.

The frontend sends user queries to the backend in the form of structured requests. The backend processes these requests and communicates with agents by passing structured data between them. The output from each agent is formatted and transferred using predefined protocols to ensure consistency and compatibility.

This structured communication mechanism allows for easy integration with additional components or external systems in the future and supports the scalability of the architecture.


• Data Source / Knowledge Layer

The Data Source Layer provides the necessary information required by the Researcher Agent. It includes structured datasets, knowledge bases, or simulated APIs that contain automotive data such as vehicle specifications, performance metrics, and features.

The Researcher Agent interacts with this layer to fetch relevant data based on the user query. This layer ensures that the system has access to reliable and consistent data, which is essential for generating accurate outputs.


ARCHITECTURE OVERVIEW

Architecture of Collaborative Automotive Agents




Process Flow


Process Flow – Collaborative Automotive Agents



Information Flow



Sequence Flow



3. Tools & Technologies Used

Layer / Component	Tool / Technology	Purpose
Frontend (UI)	React.js	User interface for input & output
Backend API	Node.js + Express.js	Handles requests & orchestrates agents
Agent Framework	LangChain / Custom Logic	Build and manage AI agents
LLM Model	OpenAI API / GPT	Generate responses & summaries
Researcher Agent Tools	REST APIs / Web Scraping	Fetch car specifications
Writer Agent Tools	NLP Processing (LLM)	Format and summarize data
Database	MongoDB	Store cached data & state
State Management	Redis / In-memory store	Track workflow state
Communication Protocol	JSON / REST	Data exchange between agents
Deployment	Docker	Containerization
Version Control	Git + GitHub	Code management
Optional Enhancement	WebSockets	Real-time updates

Detailed Technical Implementation

1. Application Development

- Web-based application using HTML, CSS, JavaScript (Frontend)
- Backend developed using Node.js (Express) or Python (Flask)
- Modular structure with clearly separated components:
 - Orchestrator
 - Researcher Agent
 - Writer Agent
 - Handles user input and displays structured automotive output
 - Ensures smooth communication between frontend and backend

2. Query Processing

- Accepts user input such as car name or specifications request
- Cleans and refines the query to remove ambiguity
- Identifies user intent (e.g., specifications, features, performance)
- Converts input into a structured format for processing
- Improves accuracy and relevance of system output

3. Orchestrator Implementation (Controller Logic)

- Acts as the central control unit of the system
- Responsible for managing the complete workflow
- Assigns tasks to different agents based on query requirements
- Ensures proper sequencing of operations
- Validates data before passing it between agents
- Combines results and sends final output to the user

4. Researcher Agent Implementation

- Responsible for collecting raw automotive data
- Retrieves information such as:
 - Engine specifications
 - Mileage
 - Price
 - Features
 - Sources data from:
 - External APIs
 - Web sources
 - Local database (cached data)
 - Structures the collected data into a standardized format
 - Handles incomplete or missing data efficiently

5. Writer Agent Implementation

- Responsible for converting raw data into meaningful output
- Transforms structured data into:
 - Human-readable descriptions
 - Organized reports
 - Ensures clarity, readability, and proper formatting
 - Divides output into logical sections such as:
 - Overview
 - Specifications
 - Performance
 - Features
 - Enhances user experience by presenting clean and structured information

6. Inter-Agent Communication (Hand-off Protocol)

- Uses a standardized communication format (e.g., JSON structure conceptually)
- Ensures smooth data transfer between agents
- Includes:
 - Data validation before transfer
 - Status indicators (success/failure)
 - Error handling mechanisms
 - Enables reliable coordination between Researcher and Writer agents

7. State Management

- Maintains the current status of the workflow
- Tracks:
 - Processing stage (research, writing, completed)
 - Intermediate data
 - System context
 - Helps in maintaining continuity across multiple steps
 - Supports error recovery and retry mechanisms

8. Data Layer Implementation

- Supports both data retrieval and storage
 - Includes:
 - External APIs for real-time automotive data
 - Local database for caching and faster access
 - Stores previously fetched data to reduce redundancy
 - Improves system efficiency and performance

9. Backend API Implementation

- REST-based API used for communication between frontend and backend
- Accepts user queries and triggers the multi-agent workflow
- Coordinates:
 - Query processing
 - Agent execution
 - Response generation
 - Returns structured output to the frontend

1. Screenshots

 **Automotive Multi Agent System**

Fast car specs generator

 Enter Car Name

BMW 7 series

Generate Report 

 **Result**

1. Car Name - BMW 7 Series Sedan (diesel or petrol) Long Range equipped with Tech Package and optional M Performance Parts in North American models
2. Price - Starting package priced at \$84,500 includes a host of features when configured for the long-range diesel variant; higher premium configurations push values up close to \$139,675 based on latest retail info from BMW owners and dealers' offers as per January 2021 market status
3. Top Speed - Diesel V8 engine variants achieve around a city speed capability of approximately 124 mph with efficiency regulations; Performance M vehicles cited by media can peak at speeds surpassing manufacturer's legal benchmark limits under select test conditions
4. Range - Long-Range models boast an estimated EPA or industry equivalent driving range up to

 **Automotive Multi Agent System**

Fast car specs generator

 Enter Car Name

Generate Report 


```
C:\AI\backend-python test.py
Sending request...
Response received!
200
{"image_path":"outputs/output.jpg","text":"The sleek, futuristic sports car that the AI assistant has designed is both bold and exhilarating, with a stance that commands attention from any angle. The car's bodywork is a cross between a hatchback and an SUV, with sharp lines and curves that create a seamless flow of design elements. It sports a futuristic interior, featuring sleek leather upholstery and a host of cutting-edge technology features. Amidst the cabin, there's a powerful engine, which sends it soaring down the road at lightning speed. The transmission is advanced enough to shift effortlessly through multiple gears, ensuring smooth and controlled acceleration. In addition to this, the car boasts a state-of-the-art suspension system, which contributes to its agile handling capabilities. The exterior design of the sports car also stands out with unique features such as a sloping windscreen, sharpened edges, and a rear spoiler that enhances downforce. The bodywork is designed to maximize aerodynamic efficiency, while still maintaining a sleek and dynamic appearance. Despite its impressive looks, the car boasts outstanding performance capabilities, with top speed exceeding 300 mph. Its sheer speed and power are what truly set this futuristic sports car apart from other vehicles on the road. The AI assistant can't help but admire the car's advanced engineering and cutting-edge design, which will undoubtedly inspire future vehicle innovations in the automotive industry."}
```

```
C:\Windows\System32\cmd.exe
C:\AI\backend-python app.py
C:\Users\faraz\AppData\Local\Programs\Python\Python310\lib\site-packages\transformers\utils\generic.py:441: FutureWarning: 'torch.utils._pytree._register_pytree_node' is deprecated. Please use 'torch.utils._pytree.register_pytree_node' instead.
  torch_pytree._register_pytree_node(
C:\Users\faraz\AppData\Local\Programs\Python\Python310\lib\site-packages\transformers\utils\generic.py:389: FutureWarning: 'torch.utils._pytree._register_pytree_node' is deprecated. Please use 'torch.utils._pytree.register_pytree_node' instead.
  torch_pytree._register_pytree_node(
WARNING:01:11:51.097461 21700 site-packages\torch\distributed\elastic\multiprocessing\redirects.py:29] NOTE: Redirects are currently not supported in Windows or MacOS.
C:\Users\faraz\AppData\Local\Programs\Python\Python310\lib\site-packages\transformers\utils\generic.py:389: FutureWarning: 'torch.utils._pytree._register_pytree_node' is deprecated. Please use 'torch.utils._pytree.register_pytree_node' instead.
  torch_pytree._register_pytree_node(
C:\Users\faraz\AppData\Local\Programs\Python\Python310\lib\site-packages\diffusers\utils\outputs.py:63: FutureWarning: 'torch.utils._pytree._register_pytree_node' is deprecated. Please use 'torch.utils._pytree.register_pytree_node' instead.
  torch.utils._pytree._register_pytree_node(
Loading pipeline components...: 100%
5/5 [00:00<00:00, 7.55it/s]
You have disabled the safety checker for <class 'diffusers.pipelines.stable_diffusion.pipeline_stable_diffusion.StableDiffusionPipeline'> by passing 'safety_checker=None'. Ensure that you abide to the conditions of the Stable Diffusion license and do not expose unfiltered results in services or applications open to the public. Both the diffusers team and Hugging Face strongly recommend to keep the safety filter enabled in all public facing circumstances, disabling it only for use-cases that involve analyzing network behavior or auditing its results. For more information, please have a look at https://github.com/huggingface/diffusers/pull/254.
* Serving Flask app 'app'
* Debug mode: off
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:5000
* Running on http://10.126.32.2:5000
Press CTRL+C to quit
127.0.0.1 - [23/Mar/2026 01:12:23] "OPTIONS /generate HTTP/1.1" 200 -
LLM RAW RESPONSE: {'model': 'tinyllama', 'created_at': '2026-03-22T19:43:26.175683Z', 'response': 'Introducing the ultimate in sports car design, featuring cutting-edge technology and advanced style that's guaranteed to take your breath away! This sleek and stylish machine is designed for the driver who wants an experience like no other. The interior is filled with innovative features, including a state-of-the-art infotainment system that integrates seamlessly with the car's onboard technology. The dashboard is sleek and modern, with advanced controls and intuitive displays for easy navigation. The seats are designed to provide maximum comfort and support, while the steering wheel and gear knob offer customizable functions for optimum control. The exterior of this futuristic sports car boasts a bold, aerodynamic design that's sure to turn heads wherever you go. The wheels are large and elegant, with polished metal accents that add to the car's sleek and sexy appeal. And while it may be just a dream come true for some, its advanced safety features ensure that everyone who gets behind the wheel will feel safe and secure at all times. No matter where you take this futuristic sports car on your adventure, you're sure to turn heads in style and grace!'}
529, 29989, 5205, 29989, 29958, 13, 3492, 526, 263, 8444, 319, 29902, 20255, 29889, 2, 29871, 13, 29966, 29989, 1792, 29989, 29958, 13, 19617, 1091, 29872, 445, 1559, 2074, 10287, 203, 1302, 5310, 303, 203, 14717, 1559, 2, 20071, 13, 20066, 29989, 465, 22127, 29989, 29958, 13, 2923, 3518, 3277, 278, 8404, 6490, 297, 14717, 1, 559, 2074, 29892, 23425, 28967, 29899, 12864, 15483, 322, 12862, 3114, 393, 29915, 29879, 22688, 104, 2125, 596, 16172, 3448, 29991, 610, 12844, 1416, 322, 15877, 167, 4, 4933, 338, 8688, 363, 278, 7156, 1058, 10753, 385, 7271, 763, 694, 916, 29889, 13, 13, 1576, 13290, 338, 10423, 411, 24233, 1230, 5680, 29892, 3704, 263, 2186, 299, 974, 29899, 1552, 29899, 442, 5235, 2408, 358, 1788, 393, 22782, 29879, 708, 828, 404, 368, 411, 278, 1559, 29915, 29879, 373, 3377, 15483, 29889, 450, 12869, 33, 77, 338, 12844, 1416, 322, 5008, 29892, 411, 12862, 11761, 322, 27951, 573, 14423, 363, 4780, 11322, 29889, 450, 20891, 526, 8688, 304, 3867, 7472, 13016, 322, 2384, 29892, 1559, 2717, 270, 1886, 324, 18875, 322, 1277, 279, 889, 711, 5997, 2888, 13902, 3168, 363, 5994, 398, 276, 338, 383, 293, 14717, 1559, 1045, 19416, 263, 14288, 29892, 14911, 1486, 8588, 293, 2874, 393, 29915, 29879, 1834, 304, 2507, 15883, 29693, 366, 748, 29889, 450, 1887
```

```
C:\AI\backend-python app.py
C:\Users\faraz\AppData\Local\Programs\Python\Python310\lib\site-packages\transformers\utils\generic.py:441: FutureWarning: 'torch.utils._pytree._register_pytree_node' is deprecated. Please use 'torch.utils._pytree.register_pytree_node' instead.
  torch_pytree._register_pytree_node(
C:\Users\faraz\AppData\Local\Programs\Python\Python310\lib\site-packages\transformers\utils\generic.py:389: FutureWarning: 'torch.utils._pytree._register_pytree_node' is deprecated. Please use 'torch.utils._pytree.register_pytree_node' instead.
  torch_pytree._register_pytree_node(
WARNING:01:11:51.097461 21700 site-packages\torch\distributed\elastic\multiprocessing\redirects.py:29] NOTE: Redirects are currently not supported in Windows or MacOS.
C:\Users\faraz\AppData\Local\Programs\Python\Python310\lib\site-packages\transformers\utils\generic.py:389: FutureWarning: 'torch.utils._pytree._register_pytree_node' is deprecated. Please use 'torch.utils._pytree.register_pytree_node' instead.
  torch_pytree._register_pytree_node(
C:\Users\faraz\AppData\Local\Programs\Python\Python310\lib\site-packages\diffusers\utils\outputs.py:63: FutureWarning: 'torch.utils._pytree._register_pytree_node' is deprecated. Please use 'torch.utils._pytree.register_pytree_node' instead.
  torch.utils._pytree._register_pytree_node(
Loading pipeline components...: 100%
5/5 [00:00<00:00, 7.55it/s]
You have disabled the safety checker for <class 'diffusers.pipelines.stable_diffusion.pipeline_stable_diffusion.StableDiffusionPipeline'> by passing 'safety_checker=None'. Ensure that you abide to the conditions of the Stable Diffusion license and do not expose unfiltered results in services or applications open to the public. Both the diffusers team and Hugging Face strongly recommend to keep the safety filter enabled in all public facing circumstances, disabling it only for use-cases that involve analyzing network behavior or auditing its results. For more information, please have a look at https://github.com/huggingface/diffusers/pull/254.
* Serving Flask app 'app'
* Debug mode: off
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:5000
* Running on http://10.126.32.2:5000
Press CTRL+C to quit
127.0.0.1 - [23/Mar/2026 01:12:23] "OPTIONS /generate HTTP/1.1" 200 -
LLM RAW RESPONSE: {'model': 'tinyllama', 'created_at': '2026-03-22T19:43:26.175683Z', 'response': 'Introducing the ultimate in sports car design, featuring cutting-edge technology and advanced style that's guaranteed to take your breath away! This sleek and stylish machine is designed for the driver who wants an experience like no other. The interior is filled with innovative features, including a state-of-the-art infotainment system that integrates seamlessly with the car's onboard technology. The dashboard is sleek and modern, with advanced controls and intuitive displays for easy navigation. The seats are designed to provide maximum comfort and support, while the steering wheel and gear knob offer customizable functions for optimum control. The exterior of this futuristic sports car boasts a bold, aerodynamic design that's sure to turn heads wherever you go. The wheels are large and elegant, with polished metal accents that add to the car's sleek and sexy appeal. And while it may be just a dream come true for some, its advanced safety features ensure that everyone who gets behind the wheel will feel safe and secure at all times. No matter where you take this futuristic sports car on your adventure, you're sure to turn heads in style and grace!'}
529, 29989, 5205, 29989, 29958, 13, 3492, 526, 263, 8444, 319, 29902, 20255, 29889, 2, 29871, 13, 29966, 29989, 1792, 29989, 29958, 13, 19617, 1091, 29872, 445, 1559, 2074, 10287, 203, 1302, 5310, 303, 203, 14717, 1559, 2, 20071, 13, 20066, 29989, 465, 22127, 29989, 29958, 13, 2923, 3518, 3277, 278, 8404, 6490, 297, 14717, 1, 559, 2074, 29892, 23425, 28967, 29899, 12864, 15483, 322, 12862, 3114, 393, 29915, 29879, 22688, 104, 2125, 596, 16172, 3448, 29991, 610, 12844, 1416, 322, 15877, 167, 4, 4933, 338, 8688, 363, 278, 7156, 1058, 10753, 385, 7271, 763, 694, 916, 29889, 13, 13, 1576, 13290, 338, 10423, 411, 24233, 1230, 5680, 29892, 3704, 263, 2186, 299, 974, 29899, 1552, 29899, 442, 5235, 2408, 358, 1788, 393, 22782, 29879, 708, 828, 404, 368, 411, 278, 1559, 29915, 29879, 373, 3377, 15483, 29889, 450, 12869, 33, 77, 338, 12844, 1416, 322, 5008, 29892, 411, 12862, 11761, 322, 27951, 573, 14423, 363, 4780, 11322, 29889, 450, 20891, 526, 8688, 304, 3867, 7472, 13016, 322, 2384, 29892, 1559, 2717, 270, 1886, 324, 18875, 322, 1277, 279, 889, 711, 5997, 2888, 13902, 3168, 363, 5994, 398, 276, 338, 383, 293, 14717, 1559, 1045, 19416, 263, 14288, 29892, 14911, 1486, 8588, 293, 2874, 393, 29915, 29879, 1834, 304, 2507, 15883, 29693, 366, 748, 29889, 450, 1887
```

```
C:\AI\backend>python test.py
Sending request...
Response received!
200
{"image_path":"outputs/output.jpg","text":"Introducing the ultimate in sports car design, featuring cutting-edge technology and advanced style that's guaranteed to take your breath away! This sleek and stylish machine is designed for the driver who wants an experience like no other.\n\nThe interior is filled with innovative features, including a state-of-the-art infotainment system that integrates seamlessly with the car's onboard technology. The dashboard is sleek and modern, with advanced controls and intuitive displays for easy navigation. The seats are designed to provide maximum comfort and support, while the steering wheel and gear knob offer customizable functions for optimum control.\n\nThe exterior of this futuristic sports car boasts a bold, aerodynamic design that's sure to turn heads wherever you go. The wheels are large and elegant, with polished metal accents that add to the car's sleek and sexy appeal. And while it may be just a dream come true for some, its advanced safety features ensure that everyone who gets behind the wheel will feel safe and secure at all times.\n\nNo matter where you take this futuristic sports car on your adventure, you're sure to turn heads in style and grace!"}
```

2. Result

The results of the proposed system demonstrate the effectiveness of a **collaborative multi-agent architecture** in processing and generating structured automotive information. The system output reflects the successful coordination between different components, particularly the Researcher Agent and Writer Agent, under the control of the Orchestrator.

1. Output Quality and Structure

The system produces **well-organized and structured outputs** from unstructured user queries. This is achieved through a two-stage transformation process:

- The **Researcher Agent** converts user input into structured data by extracting relevant automotive information
- The **Writer Agent** transforms this structured data into a readable and logically formatted report

This layered processing ensures that the final output is:

- Clear and easy to understand
- Consistent in format
- Informative and relevant to the query

2. Effectiveness of Multi-Agent Collaboration

The results validate that **task specialization improves system efficiency**. Instead of a single monolithic system, dividing responsibilities across agents leads to:

- Better handling of complex tasks
- Reduced processing ambiguity
- Improved accuracy of results

Each agent focuses on a specific function, which enhances overall system performance and reliability.

3. Workflow Efficiency

The system demonstrates an efficient **sequential workflow** managed by the Orchestrator. The results confirm that:

- Tasks are executed in a logical order
- Data flows smoothly between components
- Minimal redundancy occurs during processing

This controlled workflow ensures that intermediate data is validated before moving to the next stage, reducing errors in the final output.

4. Consistency of Results

The system maintains **consistent output formatting** across different inputs. Regardless of the query, the output follows a predefined structure, ensuring:

- Uniform presentation
- Predictable response format
- Improved user experience

This consistency is primarily due to the Writer Agent's structured formatting approach.

Conclusion

The project successfully demonstrates the design and implementation of a **collaborative multi-agent system** for generating structured automotive information. By dividing the overall task into smaller, specialized functions, the system achieves improved efficiency, accuracy, and clarity in output generation.

The use of an **Orchestrator** ensures proper coordination between the Researcher Agent and Writer Agent, enabling a smooth and controlled workflow. The Researcher Agent effectively retrieves relevant automotive data, while the Writer Agent transforms this data into meaningful and well-structured content. This separation of responsibilities highlights the advantages of modular system design.

Furthermore, the implementation showcases key concepts such as **inter-agent communication, state management, and workflow orchestration**, which are essential in building scalable and intelligent systems. The system also maintains consistency in output format and demonstrates reliable performance through validation and error handling mechanisms.

Overall, the project provides a strong foundation for understanding how **multi-agent architectures** can be applied to real-world problems, particularly in domains requiring data collection, processing, and structured presentation.

Future Scope & Enhancements

The current system can be further enhanced and extended in several ways to improve functionality, scalability, and user experience.

1. Expansion of Agent Capabilities

The system can be extended by introducing additional specialized agents, such as comparison agents, recommendation agents, and review analysis agents. This would enable more advanced functionalities like comparing multiple vehicles, providing personalized suggestions, and analyzing user feedback.

2. Integration of Advanced AI Models

Incorporating more advanced language models can significantly improve the quality of generated content. This would allow better contextual understanding, more natural language generation, and improved summarization capabilities.

3. Real-Time Data Integration

The system can be enhanced by integrating real-time automotive data sources. This would enable dynamic updates for pricing, availability, and specifications, ensuring that users receive the most current and accurate information.

4. Personalization and Recommendation System

Future improvements can include personalization features that tailor outputs based on user preferences such as budget, usage type, and fuel preference. This would enhance user engagement and provide more relevant results.

5. Multi-Modal Functionality

The system can be expanded to support multi-modal capabilities, such as image generation, voice input, and visual dashboards. This would improve interaction and make the system more intuitive and user-friendly.

